

MSENSIS MLE Task

Cats vs Dogs Image Classification

Foivos Ntoulas-Panagiotopoulos

November 21, 2025

1 Overview

This report describes my solution to the MSENSIS ML Engineering task. Given the limited time and lack of GPU, the focus was on a clean, reproducible pipeline rather than maximum accuracy.

2 Data and Preprocessing

The provided data consist of:

- a folder `dataset/images/` with JPEG images of cats and dogs,
- a CSV file `dataset/labels.csv` with:
 - `image_name`: image filename (e.g. `1234.jpg`),
 - `label`: class label (`Cat` or `Dog`, sometimes missing).

I used **Pandas** and **NumPy** for data handling. The main cleaning steps were:

1. Load `labels.csv`.
2. Drop rows with missing labels (`NaN/None`). An alternative would be to keep the unlabeled images, run k-NN in the pretrained model's feature space, and assign pseudo-labels only when the neighborhood is confident.
3. Normalize labels to lowercase and map them to integers (`cat` → 0, `dog` → 1).

Some images referenced in the CSV were missing or corrupted. To avoid runtime errors I added:

- `test.py`: checks that each `image_name` exists in `dataset/images/` and can be opened; prints a summary of missing/unreadable images.
- `clean_labels.py`: drops problematic rows, saves a backup, and overwrites `labels.csv` with a cleaned version.

After cleaning, the training set contains only valid (*image, label*) pairs and the loaders no longer crash.

3 Model and Training

Pretrained Model

I used the Vision Transformer model from Hugging Face:

- `google/vit-base-patch16-224-in21k`.

It is loaded via `ViTForImageClassification` and `AutoImageProcessor`. The original classification head (for ImageNet-21k) is replaced with a new linear layer with two outputs (cat/dog), reusing the pretrained feature extractor and only learning a small task-specific head.

Dataset and DataLoaders

A custom PyTorch Dataset (`CatsDogsDataset`):

- reads filenames and labels from `labels.csv`,
- applies the cleaning and label encoding described above,
- opens each image with PIL,
- uses the ViT image processor (resize, normalization) and
- returns `{pixel_values, labels}` for the Hugging Face model.

I use `torch.utils.data.random_split` to create an 80/20 train/validation split and wrap both subsets in `DataLoader` objects. Training is performed on CPU, so the batch size is kept relatively small (32) to avoid memory issues.

Training Procedure

Training is implemented in `train_custom_model()` in `main.py`. The setup is:

- optimizer: AdamW with learning rate 5×10^{-5} ,
- loss: cross-entropy (handled by `ViTForImageClassification` when labels are passed),
- metrics: running loss and accuracy for train and validation, shown with `tqdm` progress bars.

Due to compute limitations (CPU-only and a relatively large model), I trained for a small number of epochs (currently 1). The goal was to demonstrate the pipeline rather than fully converge the model.

After each epoch I evaluate on the validation set. If the validation accuracy improves, the model weights are saved to `models/cats_dogs_vit.pt`, which is later used for inference.

4 API and Web UI

The serving layer is implemented with **FastAPI** in the same `main.py` file.

At startup, the script:

1. Loads the same ViT architecture.
2. If a fine-tuned checkpoint exists at `models/cats_dogs_vit.pt`, it is loaded; otherwise, the base pretrained weights are used.
3. Sets the model to evaluation mode.

The application exposes:

- **GET /health**: health check returning `{"status": "ok"}`.
- **GET /**: minimal HTML page with a file upload form for manual testing.
- **POST /predict**: accepts a JPEG/PNG file, validates the type, preprocesses the image, runs a forward pass and returns JSON with:
 - `predicted_label` $\in \{"cat", "dog"\}$,
 - `confidence` $\in [0, 1]$.

This design supports both browser-based testing and programmatic use of the model.

5 How to Run

The repository includes a `README.md` with detailed steps. In summary:

1. Create a virtual environment and install dependencies with `pip install -r requirements.txt`.
2. Place the dataset under `dataset/images/` and `dataset/labels.csv`.
3. Optionally run `python test.py` and `python clean_labels.py` to clean the dataset.
4. Train the model with `python main.py -train`.
5. Start the API/Web UI with `python main.py` and open `http://localhost:8000/`.